

AIND – Adversarial Game Playing Agent

I implemented the Monte Carlo Tree Search algorithm.

Here's the performance, that is the % of successes against several baselines, over 20 tests:

	Random	Greedy	MiniMax
MCTS	100%	95%	80%

It clearly shows better performance than Random or Greedy as expected, even if it may happen that even a Greedy behavior can be sometimes successful.

It shows a certain advantage also with respect to MiniMax, as it makes a more effective search of the most promising moves through random sampling.